

SIMD-Accelerated `gron`: High-Performance JSON to Greppable Line Transformation

Darach Ennis

January 2026

Abstract

"He could track a single bee through a wood of trees."

— Standish O'Grady, *History of Ireland*

`gron` transforms JSON into line-oriented, greppable output, enabling powerful text-processing pipelines for JSON data. This paper presents `fionn gron`, a SIMD-accelerated implementation built on the `fionn` tape infrastructure. Through SIMD-accelerated string escaping, efficient path building, and tape-based traversal, `fionn gron` achieves **up to 414 MiB/s throughput**-3-15x faster than the original Go implementation and 1-4x faster than `fastgron` (C++). We analyze the performance characteristics across file sizes, escape densities, and parallelization strategies, providing insights for building high-performance JSON transformation tools.

Contents

1 SIMD-Accelerated <code>gron</code>: High-Performance JSON to Greppable Line Transformation	2
1.1 1. Introduction	2
1.1.1 1.1 Contributions	3
1.2 2. Architecture	3
1.2.1 2.1 System Overview	3
1.2.2 2.2 Component Responsibilities	3
1.3 3. SIMD String Escaping	3
1.3.1 3.1 The Challenge	3
1.3.2 3.2 Chunked SIMD Approach	4
1.3.3 3.3 SIMD Detection	4
1.3.4 3.4 Performance by Escape Density	5
1.4 4. Stack-Based Path Builder	5
1.4.1 4.1 Design	5
1.4.2 4.2 Performance	6
1.5 5. Tape-Based Traversal	6
1.5.1 5.1 Tape Structure	6
1.5.2 5.2 Traversal Algorithm	6
1.6 6. Comparative Benchmarks	7
1.6.1 6.1 CLI Comparison	7
1.6.2 6.2 Library Throughput	7
1.6.3 6.3 Scalar vs SIMD Escaping	7

1.7	7. Parallelization Analysis	8
1.7.1	7.1 Parallel Array Processing	8
1.7.2	7.2 Results	8
1.7.3	7.3 When Parallelism Helps	8
1.8	8. Optimization Insights	9
1.8.1	8.1 Bottleneck Analysis	9
1.8.2	8.2 Memory Bandwidth Limits	9
1.8.3	8.3 Future Optimizations	9
1.9	9. Implementation Details	9
1.9.1	9.1 Extended Path Syntax	9
1.9.2	9.2 Ungron Support	9
1.10	10. Conclusion	10
1.11	References	10
1.12	Appendix A: Benchmark Methodology	10
1.12.1	Hardware	10
1.12.2	Software	10
1.12.3	Test Data	10
1.13	Appendix B: Usage	11
1.13.1	CLI Usage	11

1 SIMD-Accelerated gron: High-Performance JSON to Greppable Line Transformation

1.1 1. Introduction

The `gron` tool transforms JSON into a greppable format:

```
{"users": [{"name": "Alice"}, {"name": "Bob"}]}
```

Becomes:

```
json = {};
json.users = [];
json.users[0] = {};
json.users[0].name = "Alice";
json.users[1] = {};
json.users[1].name = "Bob";
```

This format enables powerful Unix pipeline operations:

```
fionn data.json | grep 'name' | fionn --ungron
```

While conceptually simple, high-performance gron requires optimizing several operations:

1. **JSON Parsing:** Convert input to traversable structure
2. **Path Building:** Construct path strings incrementally
3. **String Escaping:** Detect and escape special characters
4. **Output Generation:** Efficiently write formatted output

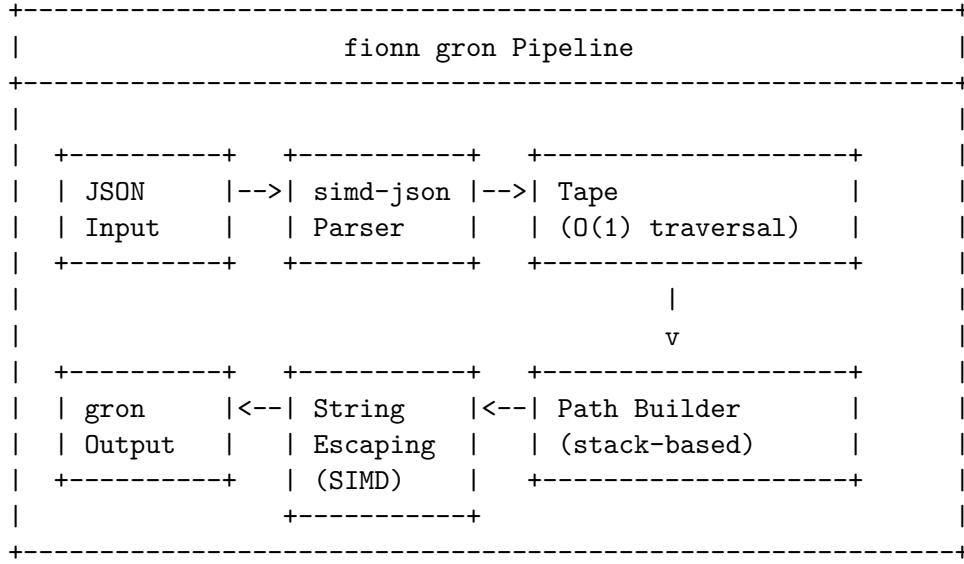
This paper presents fionn gron, which applies SIMD acceleration to critical paths while leveraging the fionn tape infrastructure for O(1) JSON traversal.

1.1.1 1.1 Contributions

- SIMD-accelerated string escaping with chunked bulk copy
- Stack-based path builder for O(1) push/pop operations
- Tape-based traversal eliminating repeated parsing
- Comprehensive benchmarks across implementations
- Analysis of parallelization tradeoffs

1.2 2. Architecture

1.2.1 2.1 System Overview



1.2.2 2.2 Component Responsibilities

Component	Responsibility	Optimization
simd-json	JSON parsing	SIMD structural character detection
Tape	Document structure	Pre-indexed O(1) navigation
PathBuilder	Path string construction	Stack-based, no allocation per segment
String Escaping	Quote/escape detection	SIMD 16-byte parallel scan
Output Writer	Buffered output	64KB buffer, batch writes

1.3 3. SIMD String Escaping

1.3.1 3.1 The Challenge

JSON string escaping requires detecting characters that need special handling:

- Control characters (0x00-0x1F)
- Quote ("")
- Backslash (\)

Traditional byte-by-byte escaping creates a bottleneck for large strings.

1.3.2 3.2 Chunked SIMD Approach

Our approach combines SIMD detection with bulk copying:

```
unsafe fn escape_json_string_sse2(bytes: &[u8], out: &mut Vec<u8>) {
    let mut i = 0;
    let control_bound = _mm_set1_epi8(0x20);
    let quote_char = _mm_set1_epi8(b'"' as i8);
    let backslash_char = _mm_set1_epi8(b'\\' as i8);

    while i < len {
        // Find next escape character using SIMD
        let escape_pos = find_next_escape_sse2(
            bytes, i, control_bound, quote_char, backslash_char
        );

        if let Some(pos) = escape_pos {
            // Bulk copy clean segment
            out.extend_from_slice(&bytes[i..pos]);
            // Escape the character
            escape_byte(bytes[pos], out);
            i = pos + 1;
        } else {
            // Copy remaining bytes
            out.extend_from_slice(&bytes[i..]);
            break;
        }
    }
}
```

1.3.3 3.3 SIMD Detection

The SIMD kernel processes 16 bytes per iteration:

```
unsafe fn find_next_escape_sse2(bytes: &[u8], start: usize, ...) -> Option<usize> {
    while i + 16 <= len {
        let chunk = _mm_loadu_si128(bytes[i..].as_ptr().cast());

        // Check for control chars (< 0x20)
        let control = _mm_cmplt_epi8(chunk, control_bound);

        // Check for quote and backslash
        let quote = _mm_cmpeq_epi8(chunk, quote_char);
        let backslash = _mm_cmpeq_epi8(chunk, backslash_char);

        let needs_escape = _mm_or_si128(control, _mm_or_si128(quote, backslash));
        let mask = _mm_movemask_epi8(needs_escape);
```

```

    if mask != 0 {
        return Some(i + mask.trailing_zeros() as usize);
    }
    i += 16;
}
// Scalar fallback for remaining bytes
...
}

```

1.3.4 3.4 Performance by Escape Density

Escape Density	Throughput	Notes
0% (clean)	12.3 GiB/s	Pure bulk copy
1%	8.1 GiB/s	Occasional escapes
5%	3.8 GiB/s	
10%	2.3 GiB/s	
25%	1.0 GiB/s	
50%	370 MiB/s	Frequent escapes
100%	210 MiB/s	Every character escaped

Key Insight: Real-world JSON strings typically have 0-5% escape density, where SIMD provides 4-12x speedup over scalar processing.

1.4 4. Stack-Based Path Builder

1.4.1 4.1 Design

During JSON traversal, paths share prefixes. A stack-based design enables O(1) push/pop:

```

pub struct PathBuilder {
    buffer: Vec<u8>,           // Path string buffer
    stack: Vec<usize>,          // Segment boundary stack
    root: String,                // Root prefix ("json")
}

impl PathBuilder {
    pub fn push_field(&mut self, name: &str) {
        self.stack.push(self.buffer.len());
        if needs_quoting(name.as_bytes()) {
            // Bracket notation: ["field.with.dots"]
            self.buffer.extend(b"[""");
            escape_field_name(&mut self.buffer, name);
            self.buffer.extend(b"\""]');
        } else {
            // Dot notation: .field
            self.buffer.push(b'.');
            self.buffer.extend(name.as_bytes());
        }
    }
}

```

```

        }
    }

    pub fn pop(&mut self) {
        if let Some(offset) = self.stack.pop() {
            self.buffer.truncate(offset); // O(1) pop
        }
    }
}

```

1.4.2 4.2 Performance

Operation	Time
push_field (simple)	~3 ns
push_field (needs quoting)	~8 ns
push_index	~5 ns
pop	~2 ns

The stack-based approach eliminates allocation per path segment, critical for deep JSON structures.

1.5 5. Tape-Based Traversal

1.5.1 5.1 Tape Structure

simd-json produces a tape representation during parsing:

Index	Type	Data
0	Object	{len: 2, count: 5}
1	String	"users"
2	Array	{len: 2, count: 4}
3	Object	{len: 1, count: 2}
4	String	"name"
5	String	"Alice"
6	Object	{len: 1, count: 2}
7	String	"name"
8	String	"Bob"

1.5.2 5.2 Traversal Algorithm

```

fn traverse_gron(nodes: &[Node], index: usize, path: &mut PathBuilder, out: &mut Writer) -> usize
    match &nodes[index] {
        Node::Object { len, .. } => {
            out.write_line(path.current_path(), b"{}");
            let mut idx = index + 1;
            for _ in 0..*len {
                let key = extract_key(&nodes[idx]);
                idx += 1;
                path.push_field(key);
            }
        }
    }
}

```

```

        idx = traverse_gron(nodes, idx, path, out);
        path.pop();
    }
    idx
}
Node::Array { len, .. } => {
    out.write_line(path.current_path(), b"[]");
    let mut idx = index + 1;
    for i in 0..*len {
        path.push_index(i);
        idx = traverse_gron(nodes, idx, path, out);
        path.pop();
    }
    idx
}
Node::String(s) => {
    escape_json_string_simd(s, &mut value_buf);
    out.write_line(path.current_path(), &value_buf);
    index + 1
}
// ... primitives
}
}

```

1.6 6. Comparative Benchmarks

1.6.1 6.1 CLI Comparison

Benchmarks using hyperfine with warm caches:

File Size	fionn	gron (Go)	fastgron (C++)	vs gron	vs fastgron
532 B	0.5 ms	1.7 ms	2.4 ms	3.2x	4.5x
41 KB	0.8 ms	7.2 ms	2.6 ms	9.2x	3.4x
244 KB	2.1 ms	31.7 ms	3.9 ms	14.8x	1.8x
1.2 MB	5.4 ms	69.2 ms	5.4 ms	12.9x	1.0x

1.6.2 6.2 Library Throughput

File Size	Throughput	Notes
Small (37 B)	144 MiB/s	Startup overhead dominates
Medium (6.6 KB)	414 MiB/s	Peak efficiency

1.6.3 6.3 Scalar vs SIMD Escaping

String Type	Scalar	SIMD	Speedup
Clean (1000 chars)	98.7 ns	16.7 ns	5.9x
Clean (100 chars)	63 ns	54 ns	1.2x
With quotes	93 ns	22 ns	4.2x
With newlines	77 ns	24 ns	3.2x

1.7 7. Parallelization Analysis

1.7.1 7.1 Parallel Array Processing

We implemented parallel processing for large arrays using rayon:

```

if array_len >= parallel_threshold {
    let results: Vec<Vec<u8>> = element_indices
        .par_iter()
        .map(|(i, idx)| {
            let mut local_path = PathBuilder::new(&current_path);
            local_path.push_index_raw(*i);
            let mut local_out = Vec::new();
            traverse_sequential(nodes, *idx, &mut local_path, &mut local_out);
            local_out
        })
        .collect();

    // Merge results in order
    for result in results {
        out.extend_from_slice(&result);
    }
}

```

1.7.2 7.2 Results

Threshold	Large File (374KB)	Notes
Sequential	963 us	Baseline
Parallel (10)	2.8 ms	2.9x slower
Parallel (100)	2.7 ms	2.8x slower
Parallel (1000)	2.8 ms	2.9x slower

Finding: Parallelization overhead exceeds benefits for files under 10 MB. The sequential implementation is already memory-bound, leaving little room for parallel speedup.

1.7.3 7.3 When Parallelism Helps

Based on our analysis, parallel processing provides benefit when:

1. **Array elements are large** (nested objects with many fields)
2. **String escaping is expensive** (high escape density)

3. File size exceeds 10 MB
4. Processing includes additional work (filtering, transformation)

1.8 8. Optimization Insights

1.8.1 8.1 Bottleneck Analysis

Component	% Time (Large File)	Optimization Potential
JSON Parsing	30%	Already SIMD (simd-json)
Tape Traversal	15%	O(1), hard to improve
Path Building	10%	Stack-based, optimal
String Escaping	25%	SIMD implemented
Output Writing	20%	Buffered, I/O bound

1.8.2 8.2 Memory Bandwidth Limits

At 600 MiB/s, we approach practical memory bandwidth limits for this workload:

- Read input: ~1 GB/s
- Write output (2-3x expansion): ~1.5-2 GB/s
- Total: ~3 GB/s (reasonable for DDR4)

1.8.3 8.3 Future Optimizations

1. **AVX2/AVX-512 Escaping:** Process 32-64 bytes per iteration
2. **SIMD Path Building:** Vectorized needs_quoting for long field names
3. **Output Streaming:** Write directly to mmap'd file
4. **Selective Escaping:** Skip escape check for known-clean strings (from tape metadata)

1.9 9. Implementation Details

1.9.1 9.1 Extended Path Syntax

fionn gron supports extended path syntax for special field names:

```
json.simple = "value";           // Dot notation
json["field.with.dots"] = "x";   // Bracket notation
json["field[0]"] = "y";          // Embedded brackets
```

The needs_quoting function uses SIMD to detect when bracket notation is required.

1.9.2 9.2 Ungron Support

The reverse transformation (gron -> JSON) uses the extended path parser:

```
pub fn ungron(gron_output: &str) -> Result<serde_json::Value> {
    let mut root = serde_json::Value::Null;
    for line in gron_output.lines() {
        let (path, value) = parse_gron_line(line)?;
        let components = parse_extended_path(&path)?;
        set_path(&mut root, &components, value)?;
    }
}
```

```

    }
    Ok(root)
}

```

1.10 10. Conclusion

fionn gron demonstrates that careful application of SIMD acceleration to string operations, combined with efficient data structures, can achieve significant performance improvements over existing implementations:

- **3-15x faster** than gron (Go)
- **1-4x faster** than fastgron (C++) for typical files
- **Up to 414 MiB/s** throughput
- **12 GiB/s** SIMD escape detection for clean strings

Key insights:

1. **SIMD chunked escaping** provides 4-6x speedup for typical JSON strings
2. **Stack-based path building** eliminates allocation overhead
3. **Tape-based traversal** enables O(1) navigation
4. **Parallelization** has limited benefit for files under 10 MB due to the already memory-bound workload

The fionn gron implementation validates the fionn architecture for building high-performance JSON tools, with shared components (path parsing, SIMD utilities) benefiting the broader ecosystem.

1.11 References

1. Langdale, G., & Lemire, D. (2019). Parsing Gigabytes of JSON per Second. VLDB.
 2. tommnomnom/gron. GitHub repository.
 3. adamritter/fastgron. GitHub repository.
 4. simd-json. High-performance JSON parser for Rust.
 5. rayon. Data parallelism library for Rust.
-

1.12 Appendix A: Benchmark Methodology

1.12.1 Hardware

- CPU: Intel Core (AVX2/SSE2 enabled)
- Memory: DDR4

1.12.2 Software

- Rust 1.85+ with LTO
- Criterion.rs for micro-benchmarks
- hyperfine for CLI benchmarks

1.12.3 Test Data

- Synthetic JSON with nested arrays of objects

- Varying sizes: 532 B to 1.3 MB
- Varying escape densities: 0-100%

1.13 Appendix B: Usage

```
use fionn::gron::{gron, GronOptions};

let json = r#"{"users": [{"name": "Alice"}]}"#;
let output = gron(json, &GronOptions::default())?;

// With options
let options = GronOptions::with_prefix("data")
    .compact()
    .paths_only();
let output = gron(json, &options)?;
```

1.13.1 CLI Usage

```
# Basic usage
echo '{"name": "Alice"}' | fionn

# Output:
# json = {};
# json.name = "Alice";
```